

Introduction to Java threads

ICS491 - Spring 2007
Concurrent and High-Performance
Programming

Henri Casanova (henric@hawaii.edu)

Threads in Programming Languages

- Several programming languages have provided constructs/abstractions for writing concurrent programs
 - Modula, Ada, etc.
- Java does it like it does everything else, by providing a Thread class
 - Note that in this class we use **J2SE 1.5.0**
- You create a thread object
- Then you can start the thread

Extending the Thread class

- To create a thread, you can extend the thread class and override its “run()” method

```
class MyThread extends Thread {  
    public void run() {  
        ...  
    }  
    ...  
}
```

```
myThread t = new MyThread();
```

Example

```
public class MyThread extends Thread {  
    public void run() {  
        for (int i=0; i<10; i++) {  
            System.out.println("Hello world #" + i);  
        }  
    }  
    ...  
}
```

```
myThread t = new MyThread();
```

Spawning a thread

- To launch, or **spawn**, a thread, you just call the thread's **start()** method
- **WARNING:** Don't call the **run()** method directly to launch a thread
 - If you call the **run()** method directly, then you just call some method of some object, and the method executes
 - Fine, but probably not what you want
 - The **start()** method, which you should not override, does all the thread launching
 - It launches a thread that starts its execution by calling the **run()** method

What happens

- The previous program runs as a process
 - Running inside the JVM
- In fact, the program runs as a single thread within a process
- When the **start()** method is called, the process creates a new thread
- We now have two threads
 - The "main", "original" thread
 - The newly created thread
- Both threads are running
 - The main thread doesn't do anything
 - The new thread prints messages to screen and exits
- When both threads are finished, then the process terminates

Example

```
public class MyThread extends Thread {
    public void run() {
        for (int i=0; i<5; i++) {
            System.out.println("Hello world #" + i);
        }
    }
}

public class MyProgram {
    public MyProgram() {
        MyThread t = new MyThread();
        t.start();
    }
    public static void main(String args[]) {
        MyProgram p = new MyProgram();
    }
}
```

Example

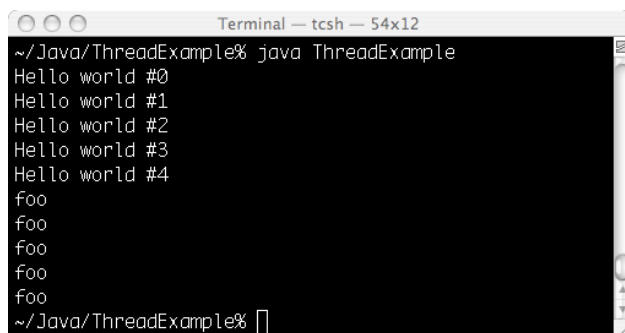
- The previous example wasn't very interesting because the main thread did nothing
 - Admittedly, this example is not interesting because the program doesn't do anything useful, but we'll get there eventually
- In fact, we could have achieved the same result with no thread at all
- So, let's have the main thread to something

Example

```
public class myThread extends Thread {
    public void run() {
        for (int i=0; i<5; i++)
            System.out.println("Hello world #" + i);
    }
}

public class MyProgram {
    public MyProgram() {
        MyThread t = new MyThread();
        t.start();
        for (int i=0; i<5; i++)
            System.out.println("foo");
    }
    public static void main(String args[]) {
        MyProgram p = new MyProgram();
    }
}
```

Example Execution



```
Terminal — tcsh — 54x12
~/Java/ThreadExample% java ThreadExample
Hello world #0
Hello world #1
Hello world #2
Hello world #3
Hello world #4
foo
foo
foo
foo
foo
~/Java/ThreadExample% 
```

- On my laptop, with my JVM, the new thread finishes executing before the main thread moves on to doing its work

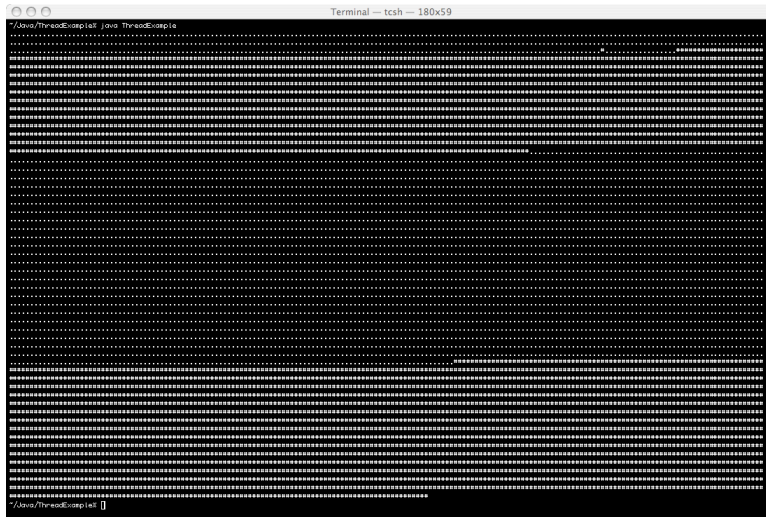
What happens?

- Now we have the main threads printing to the screen **and** the new thread printing to the screen
- Question: what will the output be?
- Answer: Impossible to tell for sure
 - If you know the implementation of the JVM on your particular machine, then you can probably tell
 - But if you write this code to be run anywhere, then you can't expect to know what happens
- Let's look at what happens on my laptop

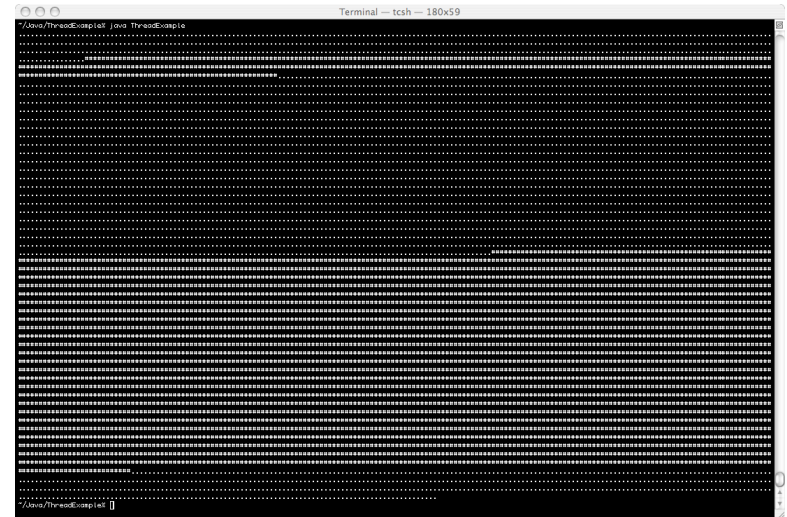
Is it really concurrent?

- One may wonder whether the execution is really concurrent
 - At least falsely concurrent
- This can be verified by having threads run for longer
- In the output that follows the new thread prints "." and the main thread prints "#"

Example Execution #1



Example Execution #2



Non-deterministic Execution

- The previous example shows what's difficult about thread programming, an especially thread debugging: it may be difficult to tell what the execution will look like
- Somebody decides when a thread runs
 - You run for a while
 - Now *you* run for a while
 - ...
- This decision process is called **scheduling**
- Let's look a little bit at the Thread class to understand this better

The Thread class

```
package java.lang;

public class Thread implements Runnable {
    public void start();
    public void run();

    public boolean isAlive();
    public Thread.State getState();

    public static void sleep(long millis);
    public static void sleep(long millis,
                             long nanos);

    // A bunch of other things we'll discuss later
    ...
}
```

The isAlive() Method

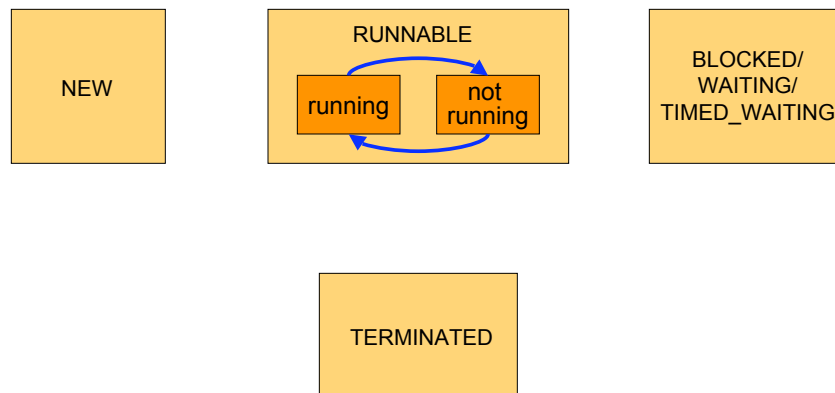
- When you spawn a thread you may not really know when or how it is going to terminate
- It may be useful to know
 - To see if the thread's work is done for instance
- The isAlive() method returns true if the thread is running, false otherwise
- Could be useful to restart a thread

```
if (!t.isAlive()) {  
    t.start();  
}
```

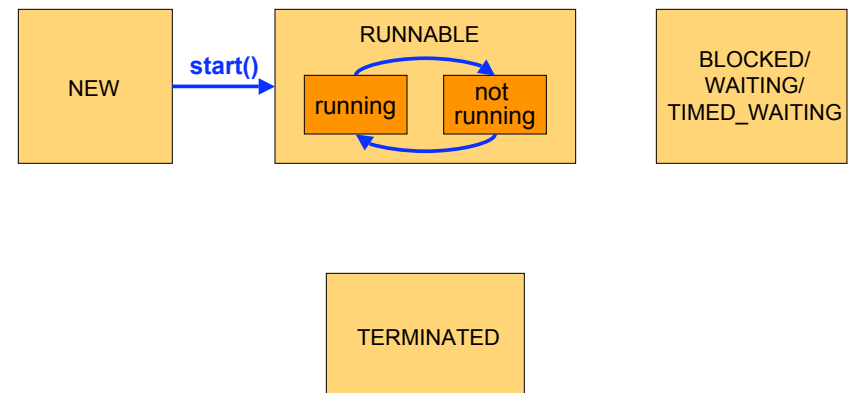
The getState() method

- The possible thread states are
 - **NEW**: A thread that hasn't been started yet
 - **RUNNABLE**: The thread can be run, and may be running as we speak
 - It may not because another runnable thread could be running
 - **BLOCKED**: The thread is blocked on a *monitor*
 - See future lecture
 - **WAITING**: The thread is waiting for another thread to do something
 - See future lecture
 - **TIMED_WAITING**: The thread is waiting for another thread to do something, but will give up after a specified time out
 - See future lecture
 - **TERMINATED**: The thread's run method has returned

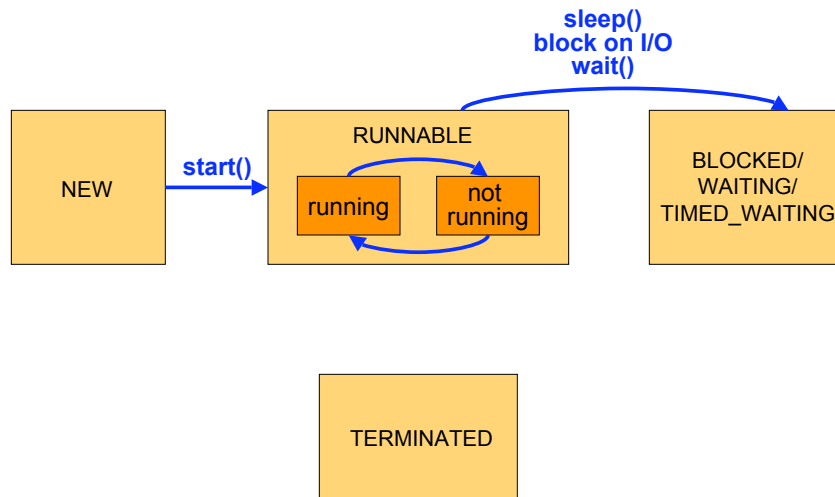
Thread Lifecycle: 4 states



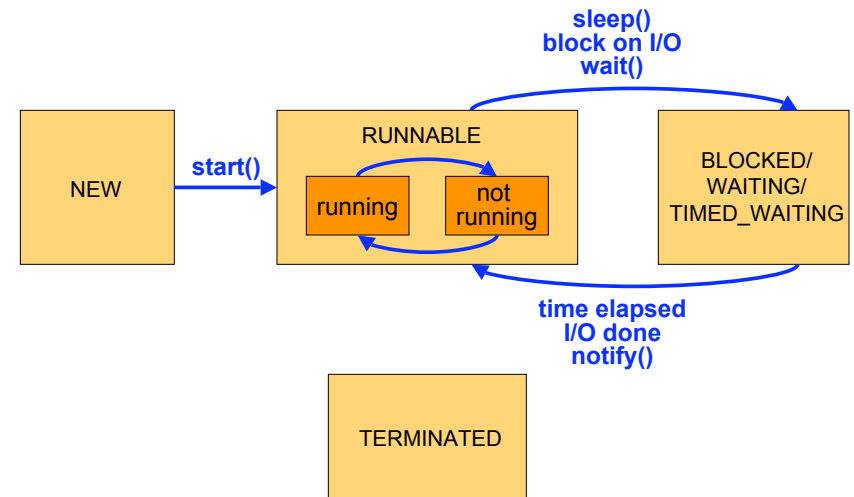
Thread Lifecycle: 4 states



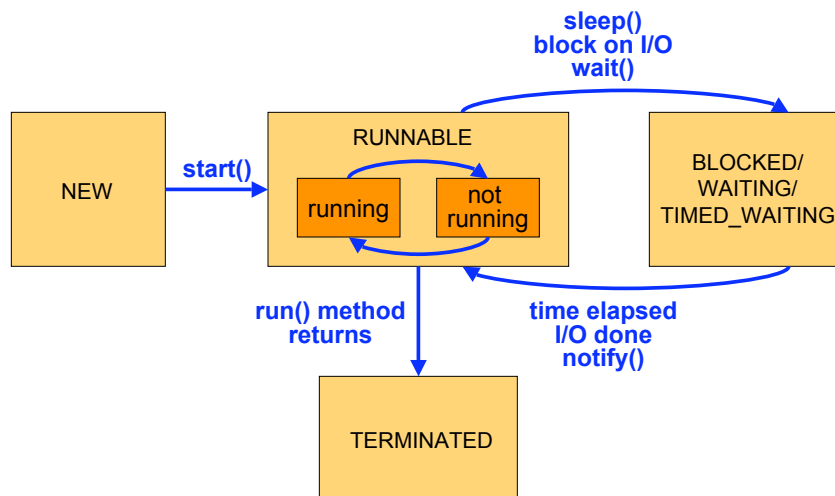
Thread Lifecycle: 4 states



Thread Lifecycle: 4 states




Thread Lifecycle: 4 states



Thread Scheduling

- The JVM keeps track of threads, enacts the thread state transition diagram
- Question: who decides which runnable thread to run?
- Old versions of the JVM used **Green Threads**
 - User-level threads implemented by the JVM
 - Invisible to the O/S




Beyond Green Threads

- Green threads have all the disadvantage of user-level threads (see previous set of lecture notes)
 - Most importantly: Cannot exploit multi-core, multi-processor architectures
- Later, the JVM provided **native threads**
 - Green threads are typically not available anymore
 - you can try to use “java -green” and see what your system says

Java Threads / Kernel Threads

- In modern JVMs, application threads are *mapped* to kernel threads



Java Threads / Kernel Threads

- This gets a bit complicated
 - The JVM has a thread scheduler for application threads, which are mapped to kernel threads
 - The O/S also has a thread scheduler for kernel threads
 - Several application threads could be mapped to the same kernel thread!
- The JVM is itself multithreaded!
- We have threads everywhere
 - Application threads in the JVM
 - Kernel threads that run application threads
 - Threads in the JVM that do some work for the JVM
- Let's look at a running JVM

A Running JVM

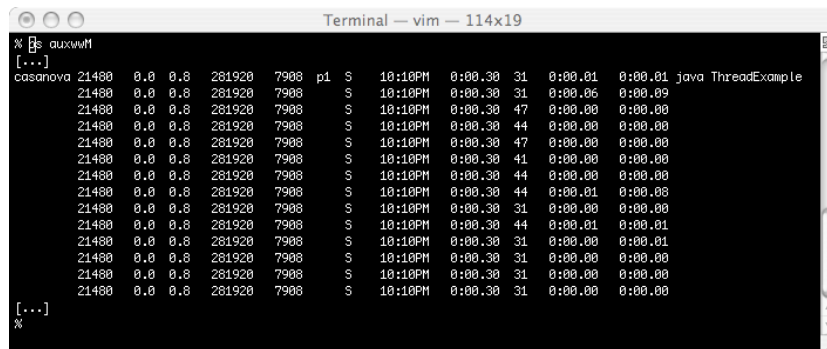
- On my laptop, a Java program that does nothing

```
Terminal — vim — 111x15
% ps auxww
...
casanova 21435  0.0  0.7  279856  7808  p3  S   10:02PM  0:00.26  31  0:00.01  0:00.01  java DoNothing
21435  0.0  0.7  279856  7808  S   10:02PM  0:00.26  31  0:00.06  0:00.09
21435  0.0  0.7  279856  7808  S   10:02PM  0:00.26  47  0:00.00  0:00.00
21435  0.0  0.7  279856  7808  S   10:02PM  0:00.26  44  0:00.00  0:00.00
21435  0.0  0.7  279856  7808  S   10:02PM  0:00.26  47  0:00.00  0:00.00
21435  0.0  0.7  279856  7808  S   10:02PM  0:00.26  41  0:00.00  0:00.00
21435  0.0  0.7  279856  7808  S   10:02PM  0:00.26  44  0:00.00  0:00.00
21435  0.0  0.7  279856  7808  S   10:02PM  0:00.26  44  0:00.01  0:00.08
21435  0.0  0.7  279856  7808  S   10:02PM  0:00.26  31  0:00.00  0:00.00
21435  0.0  0.7  279856  7808  S   10:02PM  0:00.26  44  0:00.00  0:00.00
```

- 10 threads!

A Running JVM

- On my laptop, a Java program that creates 4 threads




A terminal window titled "Terminal - vim - 114x19" showing the output of the `ps auxwwH` command. It lists 14 threads of the Java process, including the main thread and 4 application threads created by the program. Each line shows the PID, PPID, CPU%, MEM%, VSZ, RSS, TTY, STAT, TIME, and COMMAND.

- 14 threads!
 - 10 from before, one for each application thread

So what?

- At this point, it seems that we throw a bunch of threads in, and we don't really know what happens
- To some extent it's true, but we have ways to have some control
- In particular, what happens in the RUNNABLE state?



- Can we control how multiple RUNNABLE threads become running or not running?

The yield() method: example

- With the `yield()` method, a thread will pause and give other RUNNABLE threads the opportunity to execute for a while

```
public class MyThread extends Thread {
    public void run() {
        for (int i=0; i<5; i++) {
            System.out.println("Hello world #" + i);
            Thread.yield();
        }
    }
}

public class MyProgram {
    public MyProgram() {
        MyThread t = new MyThread();
        t.start();
        for (int i=0; i<5; i++) {
            System.out.println("foo");
            Thread.yield();
        }
    }
}

public static void main(String args[]) {
    MyProgram p = new MyProgram();
}
```

Example Execution



A terminal window titled "... - vim - 41x24" showing the output of the Java program. It demonstrates interleaved execution of the two threads, with "foo" and "Hello world #" messages alternating. The output is as follows:

```
% java ThreadExample
Hello world #0
foo
Hello world #1
foo
Hello world #2
foo
Hello world #3
foo
Hello world #4
foo
% java ThreadExample
Hello world #0
foo
Hello world #1
Hello world #2
foo
Hello world #3
foo
Hello world #4
foo
foo
%
```

- The use of `yield` made the threads' executions more interleaved
 - Switching between threads is more frequent
- But it's still not deterministic!
- **Programs should NEVER rely on `yield()` for correctness**
 - `yield()` is really a "hint" to the JVM

Thread Priorities

- The Thread class has a `setPriority()` and a `getPriority()` method
 - A new Thread inherits the priority of the thread that created it
- Thread priorities are integers ranging between `Thread.MIN_PRIORITY` and `Thread.MAX_PRIORITY`
 - The higher the integer, the higher the priority

So what?

- It is important to know the basics of thread scheduling to understand the behavior of concurrent programs
- One should NEVER rely on scheduling aspects to ensure correctness of the program
 - Since scheduling depends on the JVM and on the O/S, correctness due to scheduling is not portable

Thread Priorities and Scheduling

- Whenever there is a choice between multiple runnable threads, the JVM picks the higher priority one
 - High priority threads may yield to prevent **starvation** of low-priority threads
- The JVM is **preemptive**
 - If a higher priority thread is started, it gets to run
- Modern JVMs (post green threads) use **time slicing**
 - Threads of the highest priorities get chosen in a round-robin fashion
 - The use of `yield()` isn't required but, as we saw, it can increase the frequency of switching between threads
- In spite of all this:
 - The JVM can only *influence* the way in which threads are scheduled
 - Ultimately, the decision is left to the O/S

The join() method

- The `join()` method causes a thread to wait for another thread's termination
- This is useful for “dispatching” work to a worker thread and waiting for it to be done
- Let's see it used on an example

```

Terminal — vim — 70x25
public class JoinExample {

    public JoinExample() {
        Thread t1 = new Thread1();
        t1.start();
        System.out.println("Parent thread: waiting for child to finish");
        try { t1.join(); } catch (InterruptedException e) {}
        System.out.println("Parent thread: child has finished");
    }

    public static void main(String args[]) {
        JoinExample j = new JoinExample();
    }

    private class Thread1 extends Thread {
        public void run() {
            for (int i=0; i<5; i++) {
                System.out.println("Child thread: iteration"+i);
                try { Thread.sleep(1000); } catch (InterruptedException e) {}
            }
        }
    }
}

```

Runnable Example

```

public class MyTask implements Runnable {
    public void run() {
        for (int i=0; i<5; i++)
            System.out.println("Hello world #" + i);
    }
}

public class MyProgram {
    public MyProgram() {
        Thread t = new Thread(new MyTask());
        t.start();
        for (int i=0; i<5; i++)
            System.out.println("foo");
    }
    public static void main(String args[]) {
        MyProgram p = new MyProgram();
    }
}

```

The Runnable Interface

- What if you want to create a thread that extends some other class?
 - e.g., a multi-threaded applet is at the same time a Thread and an Applet
- Java does not allow for double inheritance
- Which is why it has the concept of interfaces
- So another way to create a thread is to have runnable objects
- It's actually the most common approach
 - Allows to add inheritance in a slightly easier way after the fact
- Let's see this on an example

Conclusion

- Two ways to create threads
 - extends Thread
 - implements Runnable
- Thread Scheduling is complex, not fully deterministic, and should not be counted on to guarantee program correctness