

Multi-core in JVM/Java

- Concurrent programming in java
- Prior Java 5
- Java 5 (2006)
- Java 7 (2010)
- Other topics

Basic concurrency in Java

- Java Memory Model describes how threads interact through memory
 - Single thread execution within thread as-if-serial
 - Partial order in communication between thread
- Basic concurrency construct included in language
 - Threads
 - Synchronization

Processes and Threads

- Basically, a Java virtual machine run as a single process
- Programmer can implement concurrency by using multiple threads
- It is also create a new process by instantiating `ProcessBuilder` object, e.g

```
ProcessBuilder pb = new ProcessBuilder("command", "arg1", "arg2");
```

```
Process p = pb.start();
```

Java Threads

- Provides similar features than Posix threads
- Java thread is an instance of Thread class
- Commonly used methods:

`public void run()`

`public synchronized void start()`

`public final synchronized void join(long milliseconds)`

`public static void yield()`

`public final int getPriority()`

`public final void setPriority(int newPriority)`

Java Threads

- Can be used either by subclassing *Thread* or implementing *Runnable* interface.

```
public class MyThread1 extends Thread {  
    public void run() {  
        //thread code  
    }  
    public static void main(String args[]) {  
        (new MyThread1()).start();  
    }  
}
```

```
public class MyThread2 implements Runnable {  
    public void run() {  
        // thread code  
    }  
    public static void main(String args[]) {  
        (new Thread(new MyThread2())).start();  
    }  
}
```

Synchronized Methods

- Only one thread can execute objects synchronized method(s) at time
- e.g:

```
Public class SynchronizedCounter {  
    public synchronized void update(int x) {  
        count += x;  
    }  
    public synchronized void reset {  
        count = 0;  
    }  
}
```

Synchronized Statements

- Finer-grained synchronization
- Specify the object which provides a lock

```
public class MsLunch {  
    private long c1 = 0;  
    private long c2 = 0;  
    private Object lock1 = new Object();  
    private Object lock2 = new Object();  
  
    public void inc1() {  
        synchronized(lock1){ c1++; }  
    }  
    public void inc2() {  
        synchronized(lock2) { c2++; }  
    }  
}
```

Java 5 (2006)

- **java.util.concurrent**

- Utility classes commonly useful in concurrent programming (e.g. executors, thread pools, concurrent containers)

- **java.util.concurrent.atomic**

- A small toolkit of classes that support lock-free thread-safe programming on single variables

- **java.util.concurrent.locks**

- Interfaces and classes for locking and waiting for conditions

Atomic Objects

- Package `java.util.concurrent.atomic` supports lock-free atomic operations on single variables
e.g:

```
class Sequencer {  
    private AtomicLong sequenceNumber = new AtomicLong(0);  
    public long next() {  
        return sequenceNumber.getAndIncrement();  
    }  
}
```

- Example of methods (`AtomicInteger`)

```
int addAndGet(int delta);  
boolean compareAndSet(int expect, int update);  
int decrementAndGet();  
int incrementAndGet();
```

Lock objects

- Package `java.util.concurrent.locks` provides interfaces and classes for locking and waiting for conditions
- Allow more flexibility for using locks
- Interfaces:

ReadWriteLock

Condition

Lock

Lock objects, example

```
class BoundedBuffer {  
    final Lock lock = new ReentrantLock();  
    final Condition notFull = lock.newCondition();  
    final Condition notEmpty = lock.newCondition();  
  
    final Object[] items = new Object[100];  
    int putptr, takeptr, count;  
  
    public void put(Object x) throws InterruptedException  
    {  
        lock.lock();  
        try {  
            while (count == items.length)  
                notFull.await();  
            items[putptr] = x;  
            if (++putptr == items.length) putptr = 0;  
            ++count;  
            notEmpty.signal();  
        } finally {  
            lock.unlock();  
        }  
    }  
}
```

Executor framework

- Allows to create custom thread management for *Runnable* tasks
- Decouples task submission from the mechanics of how each task will be run
- Some interfaces:
 - *Callable*
 - *Future*
 - *Executor*
 - *ExecutorService*
 - *ScheduledExecutorService*

Executor

- Examples:

```
class DirectExecutor implements Executor
{
    public void execute(Runnable r) {
        r.run();
    }
}
```

```
class ThreadPerTaskExecutor implements Executor
{
    public void execute(Runnable r) {
        new Thread(r).start();
    }
}
```

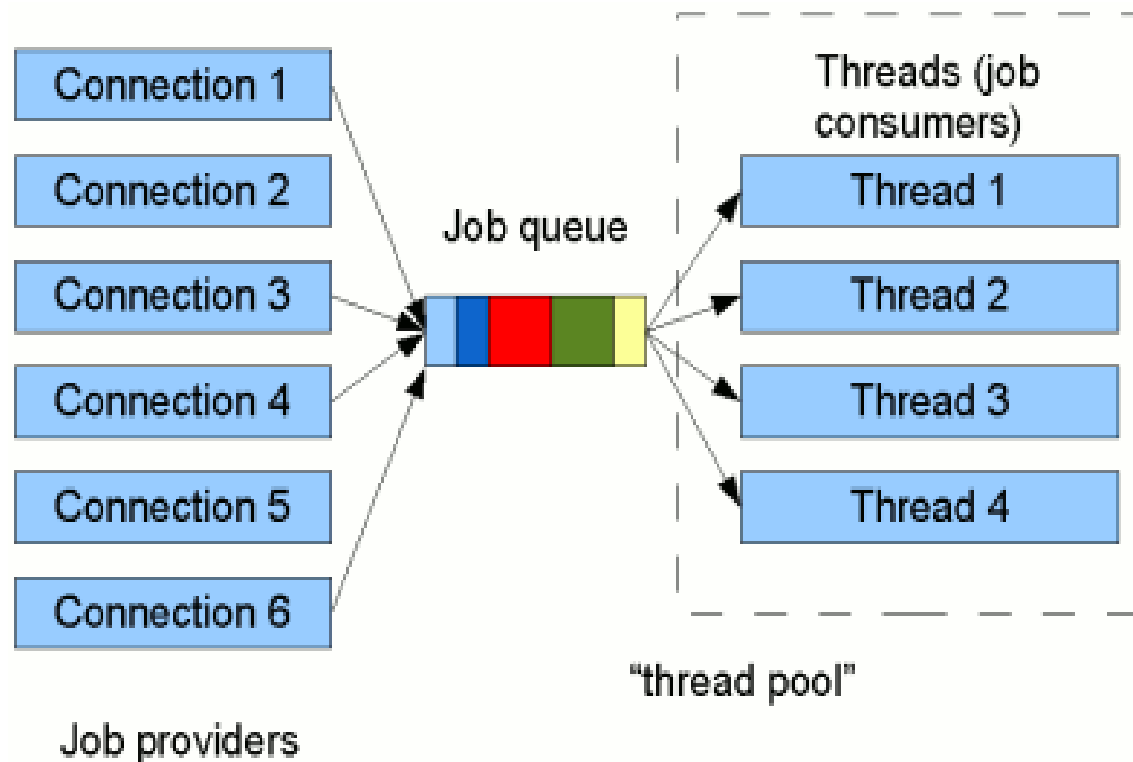
Executor

- Example of ScheduledExecutorService

```
class BeeperControl {  
    private final ScheduledExecutorService scheduler =  
        Executors.newScheduledThreadPool(1);  
  
    public void beepForAnHour() {  
        final Runnable beeper = new Runnable() {  
            public void run() { System.out.println("beep"); }  
        };  
        final ScheduledFuture<?> beeperHandle =  
            scheduler.scheduleAtFixedRate(beeper, 10, 10, SECONDS);  
        scheduler.schedule(new Runnable() {  
            public void run() { beeperHandle.cancel(true); }  
        }, 60 * 60, SECONDS);  
    }  
}
```

ThreadPoolExecutor

- Reuse threads for multiple tasks



Queues

- *ConcurrentLinkedQueue* class defines an unbounded non-blocking thread-safe FIFO
- *BlockingQueue* interface defines a thread-safe blocking queue
 - Classes: *LinkedBlockingQueue*, *ArrayBlockingQueue*, *SynchronousQueue*, *PriorityBlockingQueue*, *DelayQueue*
- *BlockingDeque* interfaces defines a thread-safe double ended queue

Synchronizers

- *Semaphore*
- *CountDownLatch*
- *CyclicBarrier*
- *Exchanger*

Concurrent Collections

- *ConcurrentHashMap*
- *CopyOnWriteArrayList*
- *CopyOnWriteArraySet*

Concurrency in Java 7

- Target release date early 2010
- Number of cores increases -> need for more and finer grained parallelism to keep processor cores busy
- Fork-join framework
- ParallelArray

Fork-join framework

- Divide and conquer approach

```
// PSEUDOCODE
Result solve(Problem problem) {
    if (problem.size < SEQUENTIAL_THRESHOLD)
        return solveSequentially(problem);
    else {
        Result left, right;
        INVOKE-IN-PARALLEL {
            left = solve(extractLeftHalf(problem));
            right = solve(extractRightHalf(problem));
        }
        return combine(left, right);
    }
}
```

Fork-join framework

- `java.util.concurrent.forkjoin`
- Designed to minimize per-task overhead
- *ForkJoinTask* is a lightweight thread
- *ForkJoinPool* hosts *ForkJoinExecutor*
- Work Stealing

Example

```
class MaxSolver extends RecursiveAction {  
    private final MaxProblem problem;  
    int result;  
  
    protected void compute() {  
        if (problem.size < THRESHOLD)  
            result = problem.solveSequentially();  
        else {  
            int m = problem.size / 2;  
            MaxSolver left, right;  
            left = new MaxSolver(problem.subproblem(0, m));  
            right = new MaxSolver(problem.subproblem(m, problem.size));  
            forkJoin(left, right);  
            result = Math.max(left.result, right.result);  
        }  
    }  
}
```

```
ForkJoinExecutor pool = new ForkJoinPool(nThreads);  
MaxSolver solver = new MaxSolver(problem);  
pool.invoke(solver);
```

Performance

Results of Running select-max on 500k-element Arrays on various systems

Threshold=	500k	50k	5k	500	50
Pentium-4 HT (2 threads)	1.0	1.07	1.02	0.82	0.2
Dual-Xeon HT (4 threads)	0.88	3.02	3.2	2.22	0.43
8-way Opteron (8 threads)	1.0	5.29	5.73	4.53	2.03
8-core Niagara (32 threads)	0.98	10.46	17.21	15.34	6.49

- Significant performance improvement can be gained if sequential threshold is reasonable
- Portable performance

ParallelArray

- Specify aggregate operation on arrays at higher abstraction layer
- Parallel Array framework automates fork-join decomposition for operation on arrays
- Supported operations:
 - Filtering
 - Mapping
 - Replacement
 - Aggregation
 - Application

ParallelArray Example

```
ParallelArray<Student> students = new ParallelArray<Student>(fjPool, data);  
double bestGpa = students.withFilter(isSenior)  
    .withMapping(selectGpa)  
    .max();
```

```
public class Student {  
    String name;  
    int graduationYear;  
    double gpa;  
}
```

```
static final Ops.Predicate<Student> isSenior = new Ops.Predicate<Student>() {  
    public boolean op(Student s) {  
        return s.graduationYear == Student.THIS_YEAR;  
    }  
};
```

```
static final Ops.ObjectToDouble<Student> selectGpa = new Ops.ObjectToDouble<Student>()  
{  
    public double op(Student student) {  
        return student.gpa;  
    }  
};
```

Parallel Array Performance

Table 1. Performance measurement for the max-GPA query
(Core 2 Quad system running Windows)

Threads	1	2	4	8
Students				
1000	1.00	0.30	0.35	1.20
10000	2.11	2.31	1.02	1.62
100000	9.99	5.28	3.63	5.53
1000000	39.34	24.67	20.94	35.11
10000000	340.25	180.28	160.21	190.41

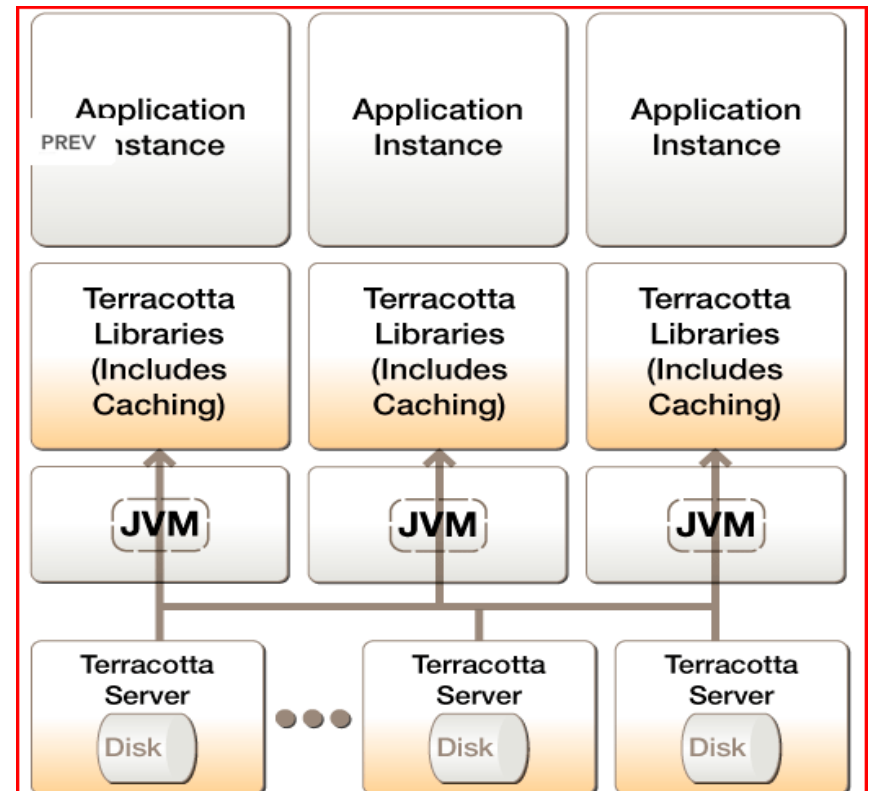
- Best speedup (2x-3x) achieved when nof cores equals to nof threads (as expected)

Other topics

- Cluster computing (Terracotta)
- Stream Programming (Pervasive DataRush)
- Highly scalable lib
- Transactional Memory

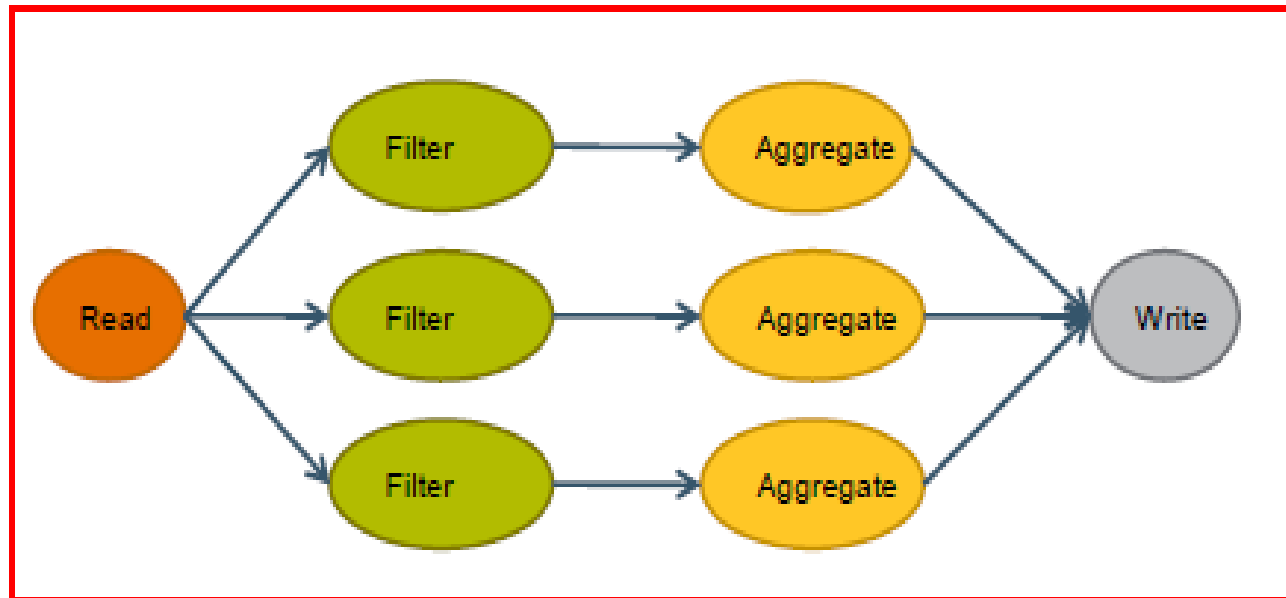
Terracotta

- Open source infrastructure to scale Java application to many computers
 - <http://www.terracotta.org>
- Transparent to programmer
 - Converts multi-threaded application to a multi-JVM (clustered) application.
 - Specify objects need to be shared across cluster



Stream Programming

- <http://www.pervasivedatarush.com/>
- Based on dataflow graph, computation nodes interconnected by queues



Highly scalable Lib

- Concurrent and Highly Scalable Collection
- <http://sourceforge.net/projects/high-scale-lib>
- Replacements for the `java.util.*` or `java.util.concurrent.*` collections

Highly Scalable Lib

- ConcurrentAutoTable
 - auto-resizing table of longs, supporting low-contention CAS operations
- NonBlockingHashMap
 - A lock-free implementation of ConcurrentHashMap
- NonBlockingSetInt
 - A lock-free bit vector set
- Linear scaling (tested up to 768 CPUs)

Transactional memory and Java

- Sequence of memory operations that execute completely (**commit**) or have no effect (**abort**)

```
atomic {  
    if (inactive.remove(p))  
        active.add(p);  
}
```

- STM or HTM
- Still a research subject

Transactional memory

Pros

- Transactions compose
- Can't acquire wrong lock
- No deadlocks
- No Priority inversion

Cons/problems

- How to roll-back I/O?
- Live-lock
- Mixing of transactional and non-transactional code
- Performance

Example (McRT STM)

