

The Design of C++0x

Bjarne Stroustrup


Texas A&M University

<http://www.research.att.com/~bs>



Overview


- Aims, Ideals, and history
- C++
- Design rules for C++0x
 - With examples
- Case study
 - Concurrency



Stroustrup - Finland 2010

8000+ Programming Languages

- C++'s family tree (part of)



- And this is a gross oversimplification!

Programming languages

- A programming language exists to help people express ideas
- Programming language features exist to serve design and programming techniques
- The primary value of a programming language is in the applications written in it
- The quest for better languages has been long and must continue



Assembler –1951

- Machine code to assembler and libraries
 - Abstraction
 - Efficiency
 - Testing
 - Documentation

THE USE OF SUB-ROUTINES IN PROGRAMMES

D. J. Wheeler

Cambridge & Illinois Universities


~~Absolutely necessary and that~~ the prime objectives
to be born in mind when constructing them are
simplicity of use, correctness of codes and accuracy
of description. All complexities should-if possible
-be buried out of sight.



1949.
May 6th
Machine in operation for first time. Printed table of
squares (0-99), time for programme 2 mins. 35 sec.
Four tanks of battery 1 in operation.

Fortran –1956

- A notation fit for humans
 - For a specific application domain
 - $A(I) = B(I) + C * D(I)$
 - Efficiency a premium
 - Portability



Simula –1967

- Organize code to “model the real world”
 - Object-oriented design
- Let the users define their own types (classes)
 - In general: map concepts/ideas into classes (“data abstraction”)
- Organize classes into hierarchies
 - Object-oriented programming



C –1974

- An simple and general notation for systems programming
 - Direct mapping of objects and basic operations to machine
 - Performance becomes somewhat portable
 - Somewhat portable



Stroustrup - Finland 2010

C with Classes –1980

- General abstraction mechanisms to cope with complexity
 - From Simula
- General close-to-hardware machine model for efficiency
 - From C
 - Became C++ in 1984
 - Commercial release 1985
 - ISO standard 1998
 - 2nd ISO standard 200x ('x' is hex ☹)




Stroustrup - Finland 2010

C++ is everywhere

- <http://www.research.att.com/~bs/applications.html>

- Telecommunications
- Google
- Microsoft applications and GUIs
- Linux tools and GUIs
- Games
- PhotoShop
- Finance
- ...



- Mars Rovers
- Marine diesel engines
- Cell phones
- Human genome project
- Micro electronics design and manufacturing
- ...

C++ ISO Standardization

- Slow, bureaucratic, democratic, formal process
 - “the worst way, except for all the rest”
 - (apologies to W. Churchill)
- About 22 nations (5 to 12 at a meeting)
- Membership have varied
 - 100 to 200+
 - 200+ members currently
 - 40 to 100 at a meeting
 - ~60 currently
- Most members work in industry
- Most members are volunteers
 - Even many of the company representatives
- Most major platform, compiler, and library vendors are represented
 - E.g., IBM, Intel, Microsoft, Sun
- End users are underrepresented



C++0X



- We have a Final Draft Standard! (as on March 13, 2010)

Design?

- Can a committee design?
 - No (at least not much)
 - Few people consider or care for the whole language
- Is C++0x designed
 - Yes
 - Well, mostly
 - You can see traces of different personalities in C++0x
- Committees
 - Discuss
 - Bring up problems
 - “Polish”



Overall goals for C++0x

- Make C++ a better language for systems programming and library building
 - Rather than providing specialized facilities for a particular sub-community (e.g. numeric computation or Windows-style application development)
 - Build directly on C++'s contributions to systems programming



- Make C++ easier to teach and learn
 - Through increased uniformity, stronger guarantees, and facilities supportive of novices (there will always be more novices than experts)

C++0x

- ‘x’ may be hex, but C++0x is not science fiction
 - Every feature is implemented somewhere
 - E.g. GCC 4.5: Rvalues, Variadic, Initializer lists, Static assertions, auto-typed variables, New function declarator syntax, Lambdas, Right angle brackets, Extern templates, Strongly-typed enums, Delegating constructors (patch), Raw string literals, Defaulted and deleted functions, Inline namespaces, Local and unnamed types as template arguments
 - Standard library components are shipping widely
 - E.g. GCC, Microsoft, Boost
 - The last design points have been settled
 - barring formal requests from National Standards Bodies
 - And they tend to be very conservative

Rules of thumb / Ideals

- Integrating features to work in combination is the key
 - And the most work
 - The whole is much more than the simple sum of its part
- Maintain stability and compatibility
- Prefer libraries to language extensions
- Prefer generality to specialization
- Support both experts and novices
- Increase type safety
- Improve performance and ability to work directly with hardware
- Make only changes that change the way people think
- Fit into the real world

Maintain stability and compatibility

- “Don’t break my code!”
 - There are billions of lines of code “out there”
 - There are millions of C++ programmers “out there”
- “Absolutely no incompatibilities” leads to ugliness
 - We introduce new keywords as needed: **auto** (recycled), **decltype**, **constexpr**, **thread_local**, **nullptr**
 - Example of incompatibility:
static_assert(4<=sizeof(int), "error: small ints");
- “Absolutely no incompatibilities” leads to absurdities
_Bool *// C99 boolean type*
typedef _Bool bool; *// C99 standard library typedef*

Support both experts and novices

- *Example*: minor syntax cleanup
vector<list<int>> vl; *// note the “missing space”*
- *Example*: simplified iteration
for (auto x : v) cout << x << '\n';
- *Note*: Experts don't easily appreciate the needs of novices
 - Example of what we couldn't get just now
string s = "12.3";
double x = lexical_cast<double>(s); *// extract value from string*

Uniform initialization

- You can use {}-initialization for all types in all contexts

```
int a[] = { 1,2,3 };
```

```
vector<int> v { 1,2,3};
```

```
vector<string> geek_heros = {  
    "Dahl", "Kernighan", "McIlroy", "Nygaard ", "Ritchie", "Stepanov"  
};
```

```
thread t{}    // default initialization
```

// remember “thread t();” is a function declaration

```
complex<double> z{1,2};    // invokes constructor
```

```
struct S { double x, y; } s {1,2}; // no constructor (just initialize members)
```

Uniform initialization

- {}-initialization **X{v}** yields the same value of **X** in every context

X x{a};

X* p = new X{a};

z = X{a}; *// use as cast*

f({a}); *// function argument (of type X)*

return {a}; *// function return value (function returning X)*

Y::Y(a) : X{a} /* ... */ *// base class initializer*

Uniform initialization

- {}-initialization does not narrow

int x1 = 7.9; *// x1 becomes 7*

int x2 {7.9}; *// error: narrowing conversion*

Table phone_numbers = {
 { "Donald Duck", 2015551234 },
 { "Mike Doonesbury", 9794566089 },
 { "Kell Dewclaw", 1123581321 }
};



Prefer libraries to language extensions

- Libraries deliver more functionality
- Libraries are immediately useful
- *Problem:* Enthusiasts prefer language features
 - see library as 2nd best
- *Example:* New library components
 - **std::thread, std::future, ...**
 - Threads ABI; not thread built-in type
 - **std::unordered_map, std::regex, ...**
 - Not built-in associative array
- *Example:* Mixed language/library extension
 - The new **for** works for every type with **std::begin()** and **std::end()**
 - The new initializer lists are based on **std::initializer_list<T>**

```
vector<string> v = { "Nygaard ", "Ritchie" };
for (auto& x : {y,z,ae,ao,aa}) cout << x << '\n';
```



Prefer generality to specialization

- *Example:* Prefer improvements to abstraction mechanisms over separate new features

- Inherited constructor

```
template<class T> class Vector : std::vector<T> {  
    using vector::vector<T>;           // inherit all constructors  
    // ...  
};
```

- Move semantics supported by rvalue references

```
template<class T> class vector {  
    // ...  
    void push_back(T&& x);           // move x into vector  
                                       // avoid copy if possible  
};
```

- *Problem:* people love small isolated features

Not a reference



Move semantics

- Often we don't want two copies, we just want to move a value

```
vector<int> make_test_equence(int n)
```

```
{
```

```
    vector<int> res;
```

```
    for (int i=0; i<n; +=i) res.push_back(rand_int());
```

```
    return res; // move, not copy
```

```
}
```

```
vector<int> seq = make_test_sequence(1000000);           // no copies
```

- New idiom for arithmetic operations:
 - **Matrix operator+(const Matrix&, const Matrix&);**
 - **a = b+c; // no copies**

Increase type safety

- Approximate the unachievable ideal
 - *Example:* Strongly-typed enumerations

```
enum class Color { red, blue, green };  
int x = Color::red;           // error: no Color->int conversion  
Color y = 7;                  // error: no int->Color conversion  
Color z = red;                // error: red not in scope  
Color c = Color::red;         // fine
```
 - *Example:* Support for general resource management
 - `std::unique_ptr` (for ownership)
 - `std::shared_ptr` (for sharing)
 - Garbage collection ABI



Move semantics in the standard library

- Prefer unique ownership to shared ownership

```
void f(shared_ptr<X>);
```

```
void g(unique_ptr<X>);
```

```
f(shared_ptr<X>(new X(1,2,3)));           // manipulate (shared) use count
```

// 0->1->2->1

```
g(unique_ptr<X>(new X(1,2,3)));           // zero overhead
```

Improve performance and the ability to work directly with hardware

- Embedded systems programming is very important
 - *Example*: address array/pointer problems
 - `array<int,7> s;` *// fixed-sized array*
 - *Example*: Generalized constant expressions (think ROM)


```
constexpr int abs(int i) { return (0<=i) ? i : -i; } // can be constant expression
```

```
struct Point {
    int x, y;
    constexpr Point(int xx, int yy) : x{xx}, y{yy} { }      // “literal type”
};
```


```
constexpr Point p1{1,2};      // must be evaluated at compile time: ok
constexpr Point p2{1,abs(x)}; // ok?: is x is a constant expression?
```

Make only changes that change the way people think

- Think/remember:
 - Object-oriented programming
 - Generic programming
 - Concurrency
 - ...
- But, most people prefer to fiddle with details
 - So there are dozens of small improvements
 - All useful somewhere
 - **long long**, **static_assert**, raw literals, **thread_local**, unicode types, ...
 - *Example*: A null pointer keyword
 - void f(int);**
 - void f(char*);**
 - f(0);** *// call f(int);*
 - f(nullptr);** *// call f(char*);*


Fit into the real world

- *Example:* Existing compilers and tools must evolve
 - Simple complete replacement is impossible
 - Tool chains are huge and expensive
 - There are more tools than you can imagine
 - C++ exists on *many* platforms
 - So the tool chain problems occur N times
 - (for each of M tools)
- *Example:* Education
 - Teachers, courses, and textbooks
 - Often mired in 1970s thinking (“C is the perfect language”)
 - or 1980s thinking (OOP: Rah! Rah!! Rah!!!)
 - “We” haven’t completely caught up with C++98!
 - “legacy code breeds more legacy code”



Areas of language change

- Machine model and concurrency Model
 - Threads library (**std::thread**)
 - Atomic ABI
 - Thread-local storage (**thread_local**)
 - Asynchronous message buffer (**std::future**)
- Support for generic programming
 - (no concepts ☹)
 - uniform initialization
 - **auto**, **decltype**, lambdas, template aliases, move semantics, variadic templates, range-**for**, ...
- Etc.
 - **static_assert**
 - improved **enums**
 - **long long**, C99 character types, etc.
 - ...



Standard Library Improvements

- New containers
 - Hash Tables (**unordered_map**, etc.)
 - Singly-linked list (**forward_list**)
 - Fixed-sized array (**array**)
- Container improvements
 - Move semantics (e.g. **push_back**)
 - Initializer-list constructors
 - Emplace operations
 - Scoped allocators
- More algorithms (just a few)
- Concurrency support
 - **thread**, **mutex**, **lock**, ...
 - **future**, **async**, ...
 - Atomic types
- Garbage collection ABI

Standard Library Improvements

- Regular Expressions (**regex**)
- General-purpose Smart Pointers (**unique_ptr**, **shared_ptr**, [...](#))
- Extensible Random Number Facility
- Enhanced Binder and function wrapper (**bind** and **function**)
- Mathematical Special Functions
- Tuple Types (**tuple**)
- Type Traits (lots)

What is C++?

**Template
meta-programming!**

A hybrid language

**A multi-paradigm
programming language**

It's C!

**Embedded systems
programming language**

**Supports
generic programming**

**A random collection
of features**

Low level!

**An object-oriented
programming language**

Stroustrup - Finland 2010



**Buffer
overflows**

Too big!

C++0x

- It *feels* like a new language
 - Compared to C++98
- How can I categorize/characterize it?
- It's *not* just “object oriented”
 - Many of the key user-defined abstractions are not objects
 - Types
 - Classifications and manipulation of types (types of types)
 - I miss “concepts”
 - Algorithms (generalized versions of computation)
 - Resources and resource lifetimes
- The pieces (language features) fit together much better than they used to

C++


**A language for
building
software
infrastructures
and resource-
constrained
applications**



**A light-weight abstraction
programming language**


Case study

- Concurrency
 - “driven by necessity”
 - More than ten years of experience



Case study: Concurrency

- What we want
 - Ease of programming
 - Writing correct concurrent code is hard
 - Portability
 - Uncompromising performance
 - System level interoperability
- We can't get everything
 - No one concurrency model is best for everything
 - De facto: we can't get all that much
 - “C++ is a systems programming language”
 - (among other things) implies serious constraints



Concurrency: `std::thread`

```
#include<thread>

void f() { std::cout << "Hello "; }

struct F { void operator()() { std::cout << "parallel world "; } };

int main()
{
    std::thread t1{f};      // f() executes in separate thread
    std::thread t2{F()};    // F()() executes in separate thread

    t1.join();             // wait for t1
    t2.join();             // wait for t2
} // spot the bug
```

Thread – pass arguments

- Use `bind()` or variadic constructor

```
void f(vector<double>&);  
struct F {  
    vector<double>& v;  
    F(vector<double>& vv) :v{vv} { }  
    void operator()();  
};  
  
int main()  
{  
    std::thread t1{std::bind(f,some_vec)};           // f(some_vec)  
    std::thread t2{f,some_vec};                     // f(some_vec)  
    t1.join(); t2.join();  
}
```

Mutual exclusion: `std::mutex`

- A **mutex** is a primitive object use for controlling access in a multi-threaded system.
- A **mutex** is a shared object (a resource)
- Simplest use:

```
std::mutex m;  
int sh; // shared data  
// ...  
m.lock();  
    // manipulate shared data:  
    sh+=1;  
m.unlock();
```



Mutex – try_lock()

- Don't wait unnecessarily

```
std::mutex m;  
int sh; // shared data  
// ...  
if (m.try_lock()) { // manipulate shared data:  
    sh+=1;  
    m.unlock();  
else {  
    // maybe do something else  
}
```


RAII for mutexes: `std::lock`

- A lock represents local ownership of a resource (the **mutex**)

```
std::mutex m;
```

```
int sh; // shared data
```

```
void f()
```

```
{
```

```
    // ...
```

```
    std::unique_lock lck(m);    // grab (acquire) the mutex
```

```
    // manipulate shared data:
```

```
    sh+=1;
```

```
}    // implicitly release the mutex
```

Potential deadlock

- Unstructured use of multiple locks is hazardous:

```
std::mutex m1;  
std::mutex m2;  
int sh1; // shared data  
int sh2;  
// ...  
void f() {  
    // ...  
    std::unique_lock lck1(m1);  
    std::unique_lock lck2(m2);  
    // manipulate shared data:  
    sh1+=sh2;  
}
```

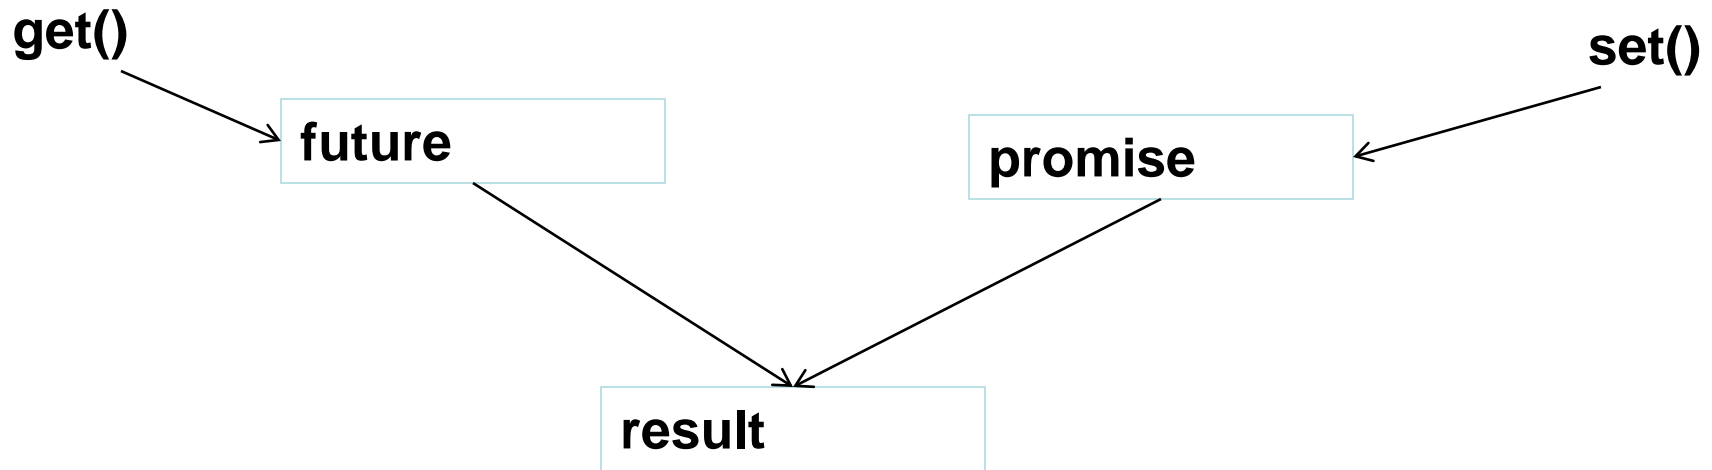


RAII for mutexes: `std::lock`

- We can safely use several locks

```
void f() {  
    // ...  
    std::unique_lock lck1(m1, std::defer_lock); // make locks but don't yet  
                                                // try to acquire the mutexes  
  
    std::unique_lock lck2(m2, std::defer_lock);  
    std::unique_lock lck3(m3, std::defer_lock);  
    // ...  
    lock(lck1, lck2, lck3);  
    // manipulate shared data  
} // implicitly release the mutexes
```

Future and promise



- future+promise provides a simple way of passing a value from one thread to another
 - No explicit synchronization
 - Exceptions can be transmitted between threads

Future and promise

- Get an **X** from a **future<X>**:
X v = f.get(); *// if necessary wait for the value to get*

- Put an **X** to a **promise<X>**:
try {
 X res;
 // compute a value for res
 p.set_value(res);
} catch (...) {
 // oops: couldn't compute res
 p.set_exception(std::current_exception());
}



async()

- Simple launcher using the variadic template interface

```
double accum(double* b, double* e, double init);
```

```
double comp(vector<double>& v) // spawn many tasks if v is large enough  
{
```

```
    if (v.size()<10000) return accum(&v[0], &v[0]+v.size(), 0.0);
```

```
    auto f0 = async(accum, &v[0], &v[v.size()/4], 0.0);
```

```
    auto f1 = async(accum, &v[v.size()/4], &v[v.size()/2], 0.0);
```

```
    auto f2 = async(accum, &v[v.size()/2], &v[v.size()*3/4], 0.0);
```

```
    auto f3 = async(accum, &v[v.size()*3/4], &v[0]+v.size(), 0.0);
```

```
    return f0.get()+f1.get()+f2.get()+f3.get();
```

```
}
```




Thanks!

- C and Simula
 - Brian Kernighan
 - Doug McIlroy
 - Kristen Nygaard
 - Dennis Ritchie
 - ...
- ISO C++ standards committee
 - Steve Clamage
 - Francis Glassborow
 - Andrew Koenig
 - Tom Plum
 - Herb Sutter
 - ...
- C++ compiler, tools, and library builders
 - Beman Dawes
 - David Vandevoorde
 - ...
- Application builders
 - Stroustrup - Finland 2010



More information

- My HOPL-II and HOPL-III papers
- The Design and Evolution of C++ (Addison Wesley 1994)
- My home pages
 - Papers, FAQs, libraries, applications, compilers, ...
 - Search for “Bjarne” or “Stroustrup”
 - “What is C++0x ?” paper
 - C++0x FAQ
- The ISO C++ standard committee’s site:
 - All documents from 1994 onwards
 - Search for “WG21”
- The Computer History Museum
 - Software preservation project’s C++ pages
 - Early compilers and documentation, etc.
 - http://www.softwarepreservation.org/projects/c_plus_plus/
 - Search for “C++ Historical Sources Archive”

